

Competitor News Daily Digest

An automated n8n workflow for AI video generation market intelligence
Nikhil Arora | August 2026

1. Workflow Overview

This workflow delivers a categorised competitor news digest by email every morning at 8:00 AM. It tracks five competitors in the AI video generation market, retrieves news published in the last 24 hours, groups articles that cover the same underlying event, filters out irrelevant matches, summarises each story, and assigns it to a fixed category taxonomy before formatting the result as an HTML email.

Pipeline stages

Stage	What happens
Schedule Trigger	Fires daily at 8:00 AM.
Google Sheets — read	Reads the competitor seed list (name, domain, aliases, tuned search query).
HTTP Request — Serper	Runs once per competitor row against Google News, restricted to the past day.
Split Out	Flattens each response's nested <code>news</code> array into individual article items.
Edit Fields	Re-attaches the competitor name to each article and normalises field names.
Remove Duplicates	Drops exact-URL duplicates across the combined multi-competitor result set.
Aggregate	Collapses all articles into a single item so the next step can compare them against each other.
AI — Clustering & relevance	Groups articles into distinct real-world stories, tags every competitor each story involves, and marks each story relevant or not.
Split Out + Filter	Unpacks stories into individual items and discards everything marked irrelevant.
AI — Summarise & categorise	Writes a 2-3 line summary per story, assigns one category, and selects the best source link.
Edit Fields	Re-attaches competitor names and stamps the run date.
Code	Groups stories by category in fixed order and builds the HTML email body.
Gmail — send	Delivers the digest.

2. Tools and APIs Used

Tool	Role	Configuration
n8n (cloud)	Workflow orchestration	Hosted n8n instance. Workflow runs on a Schedule Trigger and is exported as JSON alongside this document.
Serper.dev	News retrieval	Google News results via <code>POST https://google.serper.dev/news</code> . Authenticated with a Header Auth credential (<code>X-API-KEY</code>) stored in n8n's credential store, not in the workflow body. Free allocation of 2,500 credits; the workflow consumes 5 per run.
Google Gemini API	Clustering, relevance filtering, summarisation, categorisation	API key generated in Google AI Studio, stored as an n8n credential. Free tier; no billing attached. Two models are used: <code>gemini-3.6-flash</code> for clustering and relevance, <code>gemini-3.1-flash-lite</code> for summarisation and categorisation.
Google Sheets API	Competitor seed list	OAuth2 connection to the account owning the sheet. Read-only in this workflow.
Gmail API	Digest delivery	OAuth2 connection. Sends the digest as an HTML message.

Credentials. Every API key and OAuth token is held in n8n's credential store and referenced by name from the nodes. No key appears in the exported workflow JSON, in this document, or in any screenshot. To reproduce the workflow, create credentials with the following names and supply your own keys: Serper API Key (Header Auth), Google Gemini (API key), Google Sheets

3. Competitor Source

Competitors are stored in a Google Sheet rather than hardcoded, so the list can be edited without touching the workflow. Each row carries four fields:

Name	Domain	Aliases	Query
Runway	runwayml.com	Runway ML, RunwayML	"Runway AI" OR "RunwayML" OR "Runway Gen-4"
Kling AI	klimgai.com	Kling, Kuaishou Kling	"Kling AI" OR "Kling 3.0"
Luma AI	lumalabs.ai	Luma Labs, Luma Dream Machine, Luma Ray	"Luma AI" OR "Luma Labs"
Pika	pika.art	Pika Labs, Pika AI	"Pika AI" OR "Pika Labs"
Google Veo	deepmind.google	Veo 3, Google DeepMind Veo	"Google Veo" OR "Veo 3"

The **Query** column exists because three of the five competitor names are ordinary English words. An unqualified search for *Runway* returns fashion shows, airport runways, and startup "cash runway" articles; *Pika* returns Pokémon content. Storing a per-competitor query rather than deriving one generically lets each name be disambiguated on its own terms, which is only possible because the workflow queries competitors individually rather than in a single combined request.

4. AI Prompts Used

4.1 Clustering and relevance

Runs once per execution over the full deduplicated article set. It has to see every article at once, because deciding whether two articles describe the same event is a comparison, not a per-article judgment.

You are reviewing news articles gathered while tracking AI video generation competitors. Some are duplicate coverage of the same event from different outlets. Some mention multiple competitors in one roundup piece. Some aren't really news about any tracked competitor at all, despite matching the search.

Articles: `{{ JSON.stringify($json.articles) }}`

STEP 1 - Group into distinct real-world stories.

Two articles are the SAME story only if they describe the same specific event or announcement (e.g., the same funding round, the same launch). Similar topic or genre is NOT enough - do not merge separate comparison articles or separate listicles just because they cover similar ground.

Example - SAME story: "Kling 3.0 raises \$2.8B" + "Kling captures majority of July funding" (same funding round, different outlets).

Example - DIFFERENT stories: "Top 10 AI Video Tools" listicle + "Seedance vs Kling" comparison article (different pieces, just similar topic, even if both mention the same competitors).

For each story, list every tracked competitor it's genuinely about.

STEP 2 - Judge relevance for each story.

Mark relevant: true if the article is substantively about the tracked competitor itself - its product, technology, business, funding, partnerships, or strategy. This includes factual descriptions, reviews, and comparisons of the product, not just breaking-news events.

Mark relevant: false if any of the following apply:

- The competitor's name isn't actually mentioned (being topically about AI video in general is not enough on its own).
- The article is really about something else - a third party's content, project, or story - that merely happens to use or mention the competitor's tool in passing, rather than being about the competitor itself. Example: a meme video "made with" a competitor's tool is about the meme, not the competitor - mark false.
- The article is a generic listicle or roundup that mentions the competitor's name only as one entry among many, with no real substance about that competitor specifically.

One exception: a personnel move (an executive or employee leaving the competitor) IS relevant, even though the article's nominal subject is their new venture - it signals something real about the competitor's talent and strategy. Mark these true.

Return only a JSON object with a "stories" key containing the array.

4.2 Summarisation and categorisation

Runs once per surviving story.

You are writing one entry for a competitor-tracking digest read by a senior product manager. Be specific: name products, dollar amounts, and companies exactly. No vague adjectives like "innovative" or "game-changing".

Only use facts actually present in the title/description below. If a specific detail (dollar amount, investor name, partner name) isn't in the source text, don't invent one - write a shorter, more general summary instead of a specific-sounding but fabricated one.

Story (competitors: `{{ $json.competitors.join(', ') }}`):
`{{ JSON.stringify($json.articles) }}`

SUMMARY: 2-3 sentences, under 300 characters. Combine the most complete facts across every article - if one outlet has a number and another names a person or partner, include both. If the story involves multiple tracked competitors, make clear how each is involved.

CATEGORY: exactly one of: Product Launch, Funding, Partnership, Hiring/Leadership, Acquisition, Reviews & Comparisons, Company Update, Other.

Categorize by what the story means for the TRACKED COMPETITOR, not by the article's nominal subject. Example: an executive leaving the competitor to found a new company is Hiring/Leadership (a departure from the competitor), not Product Launch (their new venture).

Funding means the company RAISED outside investment (a round, a named

investor). Revenue, ARR, valuation mentioned without a round, or other financial metrics are Company Update, not Funding.

Reviews & Comparisons: third-party articles reviewing or comparing the product against others, with no other news content.

Other: generic descriptions or mentions with no real news value at all.

LINK: copy one URL exactly as it appears in the articles above - prefer the one with the most specific, sourced detail. Never modify or construct a URL.

Return only a JSON object.

Both prompts are paired with a **Structured Output Parser**. The categorisation schema constrains `category` to an `enum` of the eight permitted values, so an invented category fails at the schema level rather than silently reaching the digest and breaking the grouping logic.

5. Key Workflow Decisions

Per-competitor queries instead of one combined request

A single request with all five names OR'd together would have worked and used fewer API credits. It was rejected because the response would be a flat list with no indication of which competitor each article belongs to — attribution would have to be reconstructed by string-matching names back into article text, which fails precisely on the ambiguous names that needed disambiguation in the first place. Querying separately means each result set arrives already attributed, and each competitor gets its own tuned query. At five competitors per day, the credit cost of looping is immaterial.

Two distinct duplicate problems, handled by two different mechanisms

Articles duplicate in two ways. The same URL can match two competitors' searches when an article mentions both — this is exact and deterministic, so a **Remove Duplicates** node keyed on URL handles it with no AI involved. Separately, the same event gets covered by several outlets under different URLs; no string comparison catches this, so it is delegated to the clustering model, which reads titles and descriptions and judges whether two articles describe the same event. Using the cheap deterministic check first means the model only has to handle the genuinely ambiguous cases.

Relevance filtering placed at the clustering step

Keyword filtering cannot fully separate "about this competitor" from "topically adjacent." Test runs surfaced articles about an unrelated chip-materials startup and an AI-generated television series that matched the search but had nothing to do with any tracked competitor. Because the clustering step already reads every article, the relevance judgment was placed there rather than in a separate pass. This doubles as the brief's optional AI relevance-scoring extension.

Category as the digest's primary grouping key

Competitor was the initial choice, then changed. Category is the more stable axis: the number of categories is fixed at eight regardless of how many competitors are tracked, whereas grouping by competitor adds a new section header every time the list grows. The competitor name is preserved as the bolded lead-in on every line, so nothing is lost — a story is still immediately attributable at a glance.

Category taxonomy revised after seeing real output

The initial six categories (Product Launch, Funding, Partnership, Hiring/Leadership, Acquisition, Other) sent too much to *Other* on live data, and produced one clear error: an article reporting Luma AI's \$21.2M ARR was tagged *Funding*, despite the article's point being that Luma is bootstrapped and has raised nothing. Two categories were added — **Reviews & Comparisons** (recurring in this vertical, and a real market-positioning signal) and **Company Update** (revenue, headcount, and milestones, which involve dollar figures without being funding rounds). The *Funding* definition was tightened to require raised outside investment.

Grounding instruction against fabricated detail

Neither the news API nor Serper returns full article text — only a title and a short snippet. Summaries are therefore written from limited input, which creates pressure on the model to fill gaps with plausible-sounding specifics. An explicit instruction prohibits inventing amounts, investor names, or partner names not present in the source text, and directs the model to write a shorter summary instead.

Two models, split by the difficulty of the task

The Gemini free tier permits a limited number of requests per minute per model, and n8n fires per-item AI calls with no natural pause between them. Running one summarisation call per story exceeded that limit immediately on the first full test.

Rather than attach billing or slow the workflow with artificial delays, the two AI steps were split across different models according to what each actually requires. Clustering runs on **gemini-3.6-flash**: deciding whether two differently-worded articles describe the same event, and whether an article is genuinely about a tracked competitor, are comparative judgments that benefit from a more capable model. Summarisation and categorisation run on **gemini-3.1-flash-lite**, which carries a higher requests-per-minute allowance on the free tier — this is the step that fires once per story, and the task is closer to constrained formatting than reasoning, since the relevance decision has already been made upstream.

This keeps the workflow entirely within the free tier while putting model capability where it changes the output. If story volume grows enough to hit the limit again, the next step would be batching all stories into a single summarisation call rather than one call per story, or adding explicit throttling between calls.

Minimum-necessary API scope

Delivery was originally planned for Slack. The available workspace was community-owned and required admin approval for scopes the workflow did not need. Rather than request broader permissions than the task required, delivery moved to

Gmail, where the workflow authenticates as the recipient's own account and sends only to that account.

6. Sample Output

The digest below was produced by a live execution on 14 August 2026 and delivered to the configured inbox. Categories with no stories are omitted entirely rather than shown empty; on this run, Product Launch, Partnership, and Acquisition had no qualifying stories.

Competitor Digest — Aug 14

Funding

- **Kling AI** — Kling AI raised \$2.8B in funding. This investment accounted for more than half of the \$4.77B in total tech equity funding recorded in July 2026. [View]

Hiring/Leadership

- **Luma AI** — Verena Puhm, formerly Head at Luma AI, has departed the company to launch Matriarch, a new production outfit based in Los Angeles. [View]

Reviews & Comparisons

- **Kling AI** — FindArticles released a comparison of Seedance 2.5 and Kling 3.0, evaluating which AI video workflow better suits technical teams. [View]
- **Kling AI, Google Veo** — Kling 3.0 and Google Veo 3.1 were tested in a side-by-side video generation comparison against Seedance 2.0 and Sora 2. [View]
- **Google Veo** — The Hype Magazine published a comparison between Seedance 2.5 and Google's Veo 3.1, evaluating how each AI video workflow serves creators in the music and entertainment sectors. [View]

Company Update

- **Runway** — RunwayML provides an AI-powered creative platform that offers professional video editing and generation tools directly through the web. [View]
- **Luma AI** — Luma AI is currently bootstrapped with 193 employees. As of August 2026, the company's estimated ARR is \$21.2M. [View]

Two behaviours worth noting in this output. The second Reviews & Comparisons entry is tagged to both Kling AI and Google Veo — a single article covering multiple tracked competitors, correctly attributed to both rather than arbitrarily filed under one. And this run's input contained seven articles that were discarded before reaching the digest: an unrelated chip-materials startup covered by three separate outlets, two generic "best AI video tools" listicles, an AI film studio with no connection to any tracked competitor, and a meme video that merely used a competitor's tool.

7. Execution Evidence

- The workflow was executed end to end and delivered the digest above to the configured Gmail inbox on 14 August 2026.
- Serper returned results for four of five competitors; Pika had no qualifying coverage in the 24-hour window. The digest correctly omits it rather than showing an empty section.
- The clustering step merged two Singapore Business Review articles covering the same Kling funding event into a single story, and merged three separate outlets' coverage of an unrelated startup into one cluster that was then correctly discarded as irrelevant.
- Screenshots of the n8n canvas, a successful execution log, and the delivered email accompany this document.

8. Optional Extensions

Implemented — AI relevance scoring. Rather than a separate scoring pass, relevance is judged inside the clustering step, which already reads every article. Each story is returned with a boolean `relevant` flag and a Filter node discards the rejects. On the sample run this removed seven of thirteen clustered stories.

Not implemented — history logging and Friday rollout. Both were designed but not built, and are documented here as planned work.

History logging would write each story to a Notion database or Google Sheet with competitor, category, summary, source link, and run date. The design decision worth recording is that competitor and category must be stored as structured properties (Notion `Select` fields) rather than plain text. Storing them as text would make the data unpivotable later — a dashboard cannot group by a competitor name that exists only inside a sentence. Competitor should permit new options to be created automatically, since the tracked list is expected to grow; category should not, since an unexpected value there would break the digest's grouping logic and should surface as an error rather than silently expand the taxonomy.

The Friday rollout would run as a separate workflow on a weekly schedule, reading the accumulated history rather than

calling Serper again, and synthesising the week's themes and most active competitors into a single summary. It depends on history logging existing first.

9. Known Limitations

- **Summaries are written from titles and snippets, not full articles.** Neither source provides full text on a free tier. The digest is designed as a teaser with a link rather than a replacement for reading the source, but summary depth is bounded by what the snippet contains. Fetching and parsing full pages is possible but adds payroll, JavaScript-rendering, and scraping-reliability problems across dozens of outlets.
- **Grounding is instructed, not enforced.** The prompt prohibits inventing specifics, but nothing mechanically verifies a summary against its source. Summaries should be treated as pointers to be confirmed at the source link, particularly for figures and named parties.
- **Some Serper links are Google redirect URLs** rather than direct article URLs. They resolve correctly but are less readable.
- **Quiet days produce no email.** If every story is filtered out, no items reach the Gmail node and nothing sends — indistinguishable from a failed run. An IF node checking item count, with a branch sending a short "no significant competitor activity" message, would close this gap.
- **Competitor discovery is not automated.** The seed list is maintained by hand. A periodic broad category search feeding proposed additions back to the sheet was designed but not built.