



Building first-time confidence at Porter

Fifty-nine percent of first-time customers who reach the vehicle-selection screen never complete a trip. They are being asked to make one decision they have no way to make.

Nikhil Arora

Interactive prototype and a 28-page supporting-artifacts pack accompany this deck.



Success is a completed trip, not a placed booking

30 lakh

customers a month

Porter's intra-city base

10 lakh

individual customers

the segment in scope

41%

complete a first trip

measured from vehicle selection

The brief sets the target at 58% within six months. Two lines in it decide everything downstream.

The funnel opens at vehicle selection

Not at install, and not at first search. Every customer counted in the 41% has already chosen Porter and reached the screen where they must name a vehicle. Acquisition is not the problem being solved.

A trip that fails at pickup is a failure

Success is the completed trip. That rules out any answer that improves booking conversion alone — pushing more badly-sized bookings into fulfilment would move the wrong number and hurt the customer.

Every decision in this deck is tested against both lines.

Two kinds of evidence, kept apart on purpose

SECONDARY

Collected, not elicited

Review mining

427 reviews

Porter, Rapido and Uber, tagged against a single taxonomy. Nobody was asked about vehicle size, so nobody was led toward it.

Competitor desk research

6 operators

Vehicle-selection flows and sizing guidance across three markets, read directly rather than inferred from customer complaints.

One taxonomy across three apps supports a comparison no interview set could produce.

PRIMARY

21 direct conversations

Survey

17 responses

Fourteen questions, two channels, sequenced so unaided recall of pre-booking uncertainty is captured before any question names a barrier.

Interviews

4 conversations

WhatsApp, same opening prompt each time, transcripts in the pack so the questions can be judged beside the answers.

One interview is logged and excluded from every finding — the reply was not first-hand.

Five research questions drove collection.

Two were formed before any data existed. Three became answerable only once the comparison corpora were built, and are recorded that way rather than presented as predictions that happened to come true.

The loudest complaint and the real barrier are not the same thing

RANKED BY REVIEW VOLUME

share of 220 Porter reviews

Support unresponsive	20%
Price opacity	10%
Trust in an unfamiliar app	9%
Tracking visibility	6%
Damage handling	5%
Vehicle fit uncertainty	2%

RANKED BY INDEPENDENT SOURCES

survey · interview · review corpus

Vehicle fit uncertainty	3 of 3
Price opacity	3 of 3
Loading responsibility	3 of 3
Trust in an unfamiliar app	2 of 3
Damage handling	2 of 3
Support unresponsive	1 of 3

A customer who could not size their load and left never wrote a review.

The corpus is structurally blind to the barrier that stops people before they transact. Vehicle fit is the lowest-volume tag of any barrier named by all three sources — and that gap between the two rankings is the finding, not a weakness in the data.

Three first-time types, built from the responses

01

The Cautious First-Timer

Wants to try Porter and arrives with more open questions than the booking flow answers.

"Which vehicle size I needed; whether my stuff would actually fit; what the total cost would be."

Closest of the three to the conversion metric — hesitates before booking, or completes a rocky first trip.

02

The Abandoner

Opens the app, starts a booking, and walks away before the trip happens.

"Eventually I ended up not using the service."

Mid-funnel abandonment: the exact failure the brief defines success against, and the one directly observed case in the research.

03

The Offline Loyalist

Has bulky items to move, knows the apps exist, and calls a person instead — every time.

"If I had more information about the reliability and safety of app-based services, I might have considered trying one."

Has not entered the funnel at all. The competitive set here is a phone call, not another app.

A second survey wave added six responses. Every one sorted into an existing type without forcing a fourth — which is itself the useful result.

Nobody in this category answers “which vehicle?” inside the flow

Booking flows and published sizing guidance read directly across six operators in three markets, August 2026.

OPERATOR	MARKET	WHAT THE CUSTOMER GETS AT THE MOMENT OF CHOICE	WHERE SIZING HELP LIVES
Porter	India	A vehicle grid with weight limits and example loads. The catalogue itself varies by city, so the same sofa has more room to be sized correctly in Delhi than in Bangalore.	Nothing in the flow
Lalamove	Asia, Europe	Four categories described by example items — bike, car, caddy, transporter. Enter pickup and drop, then select one manually.	A blog guide, off the flow
U-Haul	United States	The most developed sizing help in the category: a room-by-room table matching home size to truck size, five self-assessment factors, and a supplies calculator.	A help article, off the flow
MOVER, Borzo, Uncle Delivery, BlackBuck	India	Positioned by use case and customer type. Where a vehicle is suggested at all, it is suggested by delivery purpose rather than by what is being moved.	Nothing

The offline operator	India	A person who asks what you are moving, and sizes the job before you commit to it.	In the conversation itself
----------------------	-------	---	----------------------------

U-Haul is the tell. A company that size does not build a room-by-room sizing guide unless customers keep getting it wrong — but it is a help article, not a step in the flow, so the work still sits with the customer at the moment they are least able to do it. The only competitor that does the translation for them is the one without an app.

On the barriers my own review corpus can compare, most weaknesses are shared — price transparency and support are weak across Porter, Rapido and Uber alike. Vehicle fit is the one with no ride-hailing equivalent, because neither app asks you to size a load.

Three barriers converged. One gates the entire funnel.

SELECTED

Vehicle fit uncertainty

The customer cannot translate a physical job — a fridge, a sofa, a room of boxes — into a vehicle category.

NOT SELECTED

Price opacity

Real and equally corroborated, but it appears identically across every competitor, and the fix crosses the pricing-engine line the brief puts out of scope.

NOT SELECTED

Loading responsibility

Sits in the same pre-booking moment and is genuinely unaddressed here. Carried forward as a named gap rather than quietly dropped.

Why vehicle fit, in priority order

1

It gates everything

Vehicle selection is mandatory and un-skippable. Every customer in the 41% passes through it; every other barrier is conditional on this step already having succeeded.

2

It fails in both worlds

The only barrier that kills the booking or kills the trip. It spans booking conversion and physical fulfilment natively, which is what the success definition demands.

3

It belongs to this category

Absent from both ride-hailing corpora, present in every goods-transport source, because neither Rapido nor Uber asks a customer to size a load.

4

Customers named the fix

Unprompted, in a question that never mentioned vehicles: a tool that scans your things and tells you which size you need.

The problem, stated

First-time customers reach the vehicle-selection screen with a concrete physical job and are asked to convert it into an abstract vehicle category. Nothing in the app tells them which one fits. Repeat and business users have built that judgement across dozens of trips; first-timers guess. The guess is where the funnel breaks — some cannot make it and leave, and some make it wrong and find out when the vehicle arrives.

01

Vehicle
selection

02

Trip
details

03

Fare
estimate

04

Booking
placed

05

Driver
assigned

06

Pickup &
loading

07

In
transit

08

Trip
completed

Leak 1 • Configuration

Stages 1 to 4 — the steepest

Cannot size the job, exits before booking. Structurally invisible to review data, because someone who leaves before transacting has nothing to review.

Leak 2 • Fulfilment

Stages 5 to 8

Sized it wrong, discovers it when the vehicle arrives. Shows up in review data only as its consequences: damage, renegotiation at the curb, cancellations.

Root cause — no step in the flow translates a physical job into a vehicle specification.

What closing the gap is worth

OPPORTUNITY SIZE

Total market

30 lakh / month

Serviceable

10 lakh / month

Obtainable

1.5 lakh / month

Every Porter intra-city customer, then individual customers only, then the first-timers who reach vehicle selection. Only the obtainable layer is claimed — nothing in the model takes credit outside it.

Stress-tested on its softest input

First-timer share is the assumption everything else rests on. At 10% the annual value is ₹28.6 Cr; at 20% it is ₹57.1 Cr. The case for building holds across the whole range — only its size moves.

VALUE OF THE MOVE, PER ANNUAL COHORT

₹42.8 Cr

Net expected value across 18 lakh first-time customers a year, on a lifetime-value horizon of 24 months.

61,500 → 87,000 completed first trips a month

PER CUSTOMER REACHING VEHICLE SELECTION

Expected value today ₹324

Expected value at target ₹562

+73% in net expected value per first-time customer

Volumes, order value and repeat rate are stated assumptions, carried consistently and marked as such throughout the pack. The brief supplies 41%, 58% and the customer base.

Ten solutions, scored on reach, impact, confidence and effort


Reach is absolute customers a month at the moment the barrier strikes. 1.50 lakh first-timers reach vehicle selection; interventions firing after a booking exists reach 0.82 lakh, and mid-flow exits 0.68 lakh. Effort is a person-week band — 1, 2 or 4 — applied identically to every row, so the arithmetic on each line is checkable.

#	SOLUTION	REACH / MO	IMPACT	CONF.	EFFORT	SCORE
1	Vehicle-sizing assistant — describe or photograph the load, get a vehicle with reasons	1.50 L	3	80%	4	0.90
2	Guided item configurator — rules-based lookup, no model	1.50 L	2	80%	1	2.40
3	Capacity confirmation before dispatch	0.82 L	2	80%	2	0.66
4	Itemised upfront fare breakdown	1.50 L	2	80%	1	2.40
5	Pre-booking orientation on packing and lifting	1.50 L	2	50%	1	1.50
6	First-time orientation content at the entry point	1.50 L	1	50%	1	0.75
7	Pre-booking route and driver transparency	1.50 L	1	50%	2	0.38
8	Save and resume on a mid-flow exit	0.68 L	2	50%	2	0.34
9	Fare lock with an in-app dispute flag	0.82 L	1	50%	2	0.21
10	Proactive trip-status updates	0.82 L	0.5	25%	2	0.05

The top two are not the top-scoring pair. That is deliberate.

BUILT Vehicle-sizing assistant 0.90 The only solution scoring maximum impact, the most direct answer to the barrier, and the only one that carries a trust layer alongside the recommendation. A customer asked for exactly this, unprompted, in a question that never mentioned vehicles.	RUNNER-UP, ARGUED NOT BUILT Guided item configurator 2.40 Tied for the highest score and answers the same barrier from the opposite direction — deterministic and instant against guided and evidenced. That opposition is what makes the guidance-versus-speed trade-off concrete instead of generic.
Why override the score Two solutions tie at the top of the table and neither answers the primary barrier as directly. A framework ranks; it does not decide. Keeping both choices aligned to one barrier is what produces a coherent product rather than a diversified list.	
The trade-off both solutions share Porter competes on on-demand speed, and both options buy certainty by adding a step to the flow. The configurator costs less time per customer; neither is free. Time from vehicle selection to booking is measured at cohort level for exactly this reason — a trade-off that is named and never measured was only ever rhetorical.	

Additive at two points, both before a driver is dispatched



Stages two, three, five, six and seven are untouched. The solution adds one flow, one conditional confirmation and one post-trip prompt, and restructures nothing.


Why the fulfilment fix sits before dispatch, not at the curb

A capacity check at pickup was designed twice and rejected twice. At the curb, with the load half-carried, a driver naming a cash number beats any in-app flow — and it depends on that driver choosing to log a mismatch he profits more from settling informally. Driver pay is out of scope to change.

What moving it upstream buys


No driver to compete with, and no dispatch to unwind, because the booking has not been created yet. A switch is a re-selection rather than a re-assignment, which keeps the design clear of backend dispatch. The cost of being wrong becomes a fare difference the customer chooses.

The flow, end to end




The entry point

One opt-in prompt above an unchanged grid




Describe the job

One input bar, not a mode choice




The recommendation

Vehicle, reasoning and evidence on one screen




When it is unsure

Questions, not a quieter guess



Before a driver is sent

The last moment a change is free



Back to the existing flow

Unchanged Porter screens

Interactive prototype • claude.ai/code/artifact/751e3234-cef2-4469-90c0-916a988222c2

Eight screens, six of them shown here. Matching runs on deterministic rules over the typed description, tested against a fixed example set.

Confidence comes from a recommendation that is explained, not asserted

A vehicle name on its own is another claim the customer has no way to evaluate — which is the position they were already in. The design spends its complexity on the two things that make a claim checkable: why this vehicle, and what happened to people who moved a similar load in it.

01

Opt-in, not a gate

The existing grid stays the default. Gating the screen would force the flow onto repeat and business customers who are out of scope, and would restructure a core screen rather than add to it.

The cost is real: a customer who does not notice the prompt gets no help. That caps the solution at its own discoverability, which is why placement and wording are tested from the first cohort.

02

Reasoning is unavoidable

The vehicle, the one-line reason and the comparable moves sit on one screen with no tabs and no disclosure.

Putting the reasoning behind a tap would have been cleaner and would have destroyed the mechanism. Nobody taps “why” on a recommendation they have no reason to doubt yet.

03

Low confidence asks

When the description is too vague to size, the system asks three structured questions instead of returning a hedged guess.

A confident wrong answer manufactures exactly the false certainty this work exists to remove — and it arrives with the app's endorsement attached.

04

Override is never blocked

Accepting pre-fills the vehicle; the full grid is one tap away at all times.

If the job is to build confidence, it has to be possible to disagree. A recommendation that cannot be dismissed reads as a decision made on the customer's behalf.

Four releases, each scoped to what it can prove

RELEASE 0

Instrumentation

Pre-launch

SHIPS

The post-trip confirmation ships to every completed first trip. No customer-facing assistant yet.

WHY HERE

The recommendation screen shows comparable past moves, and that evidence does not exist today. Shipping the instrumentation first means the first cohort opens against a populated corpus instead of an empty one.

RELEASE 1

The assistant

Cohorts 0 – 3

SHIPS

Photo and text input, the low-confidence questions, and the load check before dispatch.

WHY HERE

Photo and text ship together. The fallback is the expensive shared component and gets built either way, so staggering them defers the capability customers asked for without deferring any of the cost.

RELEASE 2

Dimensional accuracy

After general availability

SHIPS

Longest-side gating wherever load-bay dimensions are published, and a live per-city catalogue feed.

WHY HERE

Answers the two failure modes most likely to produce a confidently wrong answer. Deferred rather than dropped: both need real accuracy data before they can be calibrated against anything.

RELEASE 3

Segment expansion

Triggered, not dated

SHIPS

Load types beyond single-item and small multi-item household goods, widened one type at a time.

WHY HERE

Accuracy tracks how dense the evidence is, so scope widens where the loop has already run rather than where a roadmap says it should.

Scope is set by what each release can prove, not by what the design can hold.

Launch on a use case, not a city

Scope — binding

Full home relocation is the packers-and-movers job, which the brief puts out of scope. Launching there would put the first cohort outside the case study's own boundary, so this is not a preference: the alternative is disallowed.

Tractability

A bounded set of item types can be enumerated. Full relocations are combinatorially open — every one is a different list — which would force both the model and its fallback to be general on day one.

Evidence — directional

The loads named across the research are single items and small household sets: a car seat, a sofa, a fridge, a room of cartons. On twenty-one conversations that points rather than proves — it corroborates a segment the two arguments above select on independent grounds.

This is an in-product launch, not a marketing one.

The customers this serves are already inside the funnel by the metric's own definition, so adoption reduces to one thing: whether they notice the prompt. Levers are placement and wording, first-time flagging, and a distinct treatment for a customer's very first arrival at the screen. Geography becomes a rollout dimension with a floor, because Porter's vehicle catalogue is deeper in some cities than others.

COHORT ROLLOUT — THE SIX-MONTH TARGET SETS THE WINDOWS

Shadow

0%

WEEKS 1-3

The model runs silently on live sessions, its recommendation compared against what the customer chose and what actually completed.

Advance when shadow accuracy beats unaided manual selection. If it does not, nothing downstream matters.

Pilot

5%

WEEKS 4-13

Two or three metros chosen on first-timer volume and catalogue depth. Full flow live against a concurrent holdout.

Advance on no guardrail regression. Absolute targets are set from this cohort's own data, not invented before it.

Scale

25%

WEEKS 14-20

Wider city set, same load segment. The evidence loop is now self-sustaining.

Advance when the north star is directionally positive against the holdout and the fallback has run at real volume.

General

100%

WEEK 21 →

All cities. Load types widen one at a time as the evidence base allows. A small holdout is kept permanently.

The ninety-day repeat signal reads here first — the earliest point the repeat-rate assumption is constrained by data.

Two loops with one shared hinge

THE EVIDENCE LOOP • COMPOUNDS ACROSS CUSTOMERS

- 1 A first-timer describes an unfamiliar load
- 2 The assistant recommends, and shows comparable past moves
- 3 Confidence to book — the trip completes correctly
- 4 The post-trip prompt captures what moved and whether it fitted
- 5 That trip becomes the evidence shown to the next customer

Each completed trip makes the next recommendation better evidenced, for a different customer. This is why the solution gets more valuable at scale rather than merely bigger.

THE TRUST CHAIN • RUNS ONCE, PER CUSTOMER


- 1 A correct first trip
- 2 “The app got the vehicle right”
- 3 The customer now has their own sizing judgement
- 4 Second booking — faster, and without the assistant
- 5 Repeat customer, who has left the first-time segment

The customer graduates, which is the intended end state. Assistant engagement will therefore never compound — and its flatness must not be read as decline. Health is read on cohort conversion, never on daily active use.


Both loops turn on the same event: a trip that completes correctly.

A badly-sized first trip does three things at once — it fails the customer, it produces no evidence for the next one, and it produces no repeat customer. One hinge, three losses. That is why the design protects it upstream, at the last moment a change still costs nothing, rather than at the curb where it costs a negotiation.

One north star, leading and lagging indicators, five guardrails



GUARDRAILS — EVERY ONE GATED AGAINST A CONCURRENT HOLDOUT



Below the north star these are holdout gates, not invented targets — and the two diagnostics carry decision rules that say which fix, not just how often.

Seventeen risks. Three decide whether this works.

Described wrongly, sized confidently

A load the customer describes inaccurately produces a confident recommendation, passes every check, and fails at the curb. There is no mitigation inside this version.

So it is measured rather than assumed away: the dispatch record is a floor, and the customer's own post-trip answer catches what a driver settling in cash will never log. The gap between the two is itself a finding.

Weight is the easy half

Five vehicle categories sit within a few hundred kilograms of each other, and load-bay dimensions are published for only some. A sofa is a length problem wearing weight's clothing.

The oversized-item question becomes mandatory whenever a single large item is named, the recommendation gates on longest side wherever a dimension exists, and it rounds up where one does not.

A correct answer can still feel like a penalty

Getting the vehicle right frequently means a bigger, costlier vehicle, and this solution deliberately does not touch price transparency.

Fare disputes is therefore the guardrail most likely to move, and pricing changes are out of scope. The only lever left is language: naming what the cheaper option actually costs turns an increase into a purchase.

WHAT THIS DESIGN CONSCIOUSLY DOES NOT OPTIMISE FOR

Reach

opt-in, so the ceiling is discoverability

Recovery

prevention before dispatch, nothing at the curb

The cheapest trip

the trip that completes instead

Barrier breadth

depth on one instead of gestures at four

Revenue per trip

some current revenue comes from mis-sizing

Kill criteria are set in advance, so the decision is not made under pressure with the data already in the room: accuracy below the manual baseline halts the rollout; a material regression in fare disputes rolls it back a cohort.

One barrier, chosen from fourteen. One solution, built and measurable.

₹42.8 Cr

annual value of moving first-trip completion from 41% to 58%

3 of 3

independent sources naming vehicle fit — across four markets and four services

2 points

where the flow changes, both of them before a driver is dispatched

Interactive prototype

claude.ai/code/artifact/751e3234-cef2-4469-90c0-916a988222c2

Supporting artifacts

28 pages — research base, corpora, personas, competitive analysis, requirements, metrics and the full risk register